

CONTAINERS AND KUBERNETES WITH NSX

Anthony Burke

Solutions Architect

Network and Security Business Unit,

Customer Success | Ambassador, Office of the CTO

@pandom_

POSSIBLE
BEGINS
WITH YOU



AGENDA

- Containers
- Kubernetes
- Kubernetes and NSX-T
- Planespotters

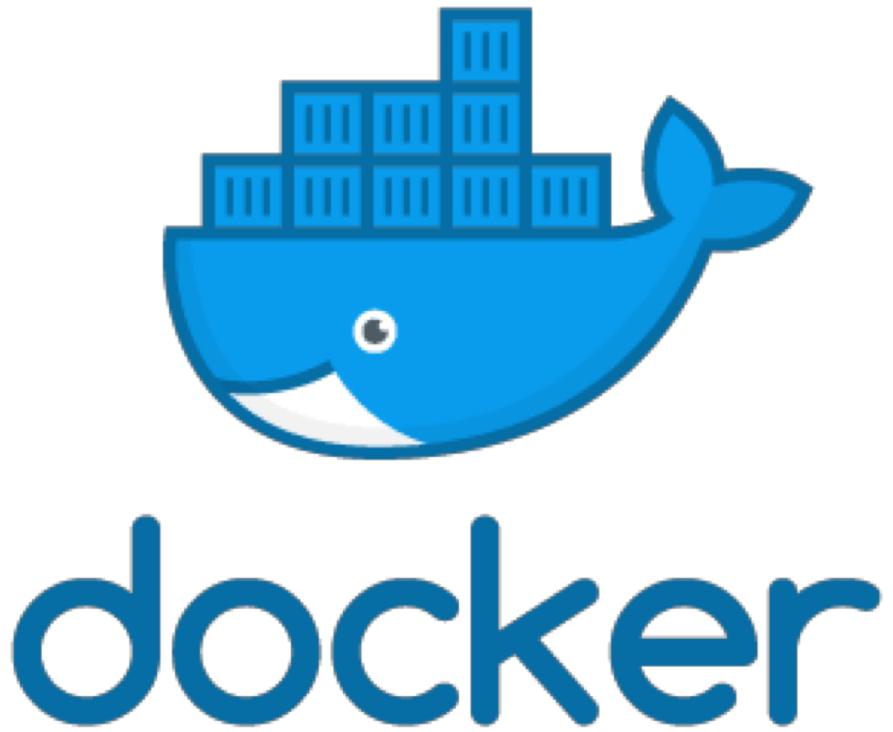
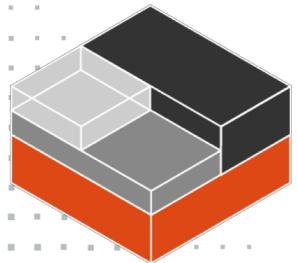


Containers

What is a container?

We're running containers

A container runtime



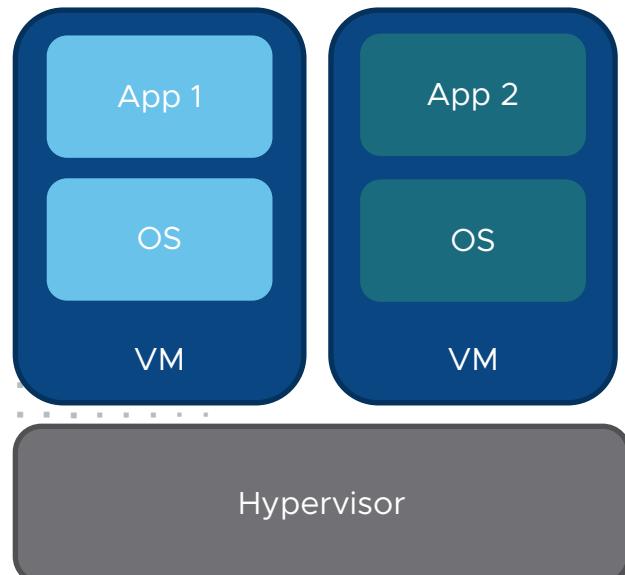
How containers are “sold” – made human friendly

Virtual Machine can run one or many apps

Libraries and executables are shared

Consume the resources allocated to VM

Each VM has a full OS / App / Binaries

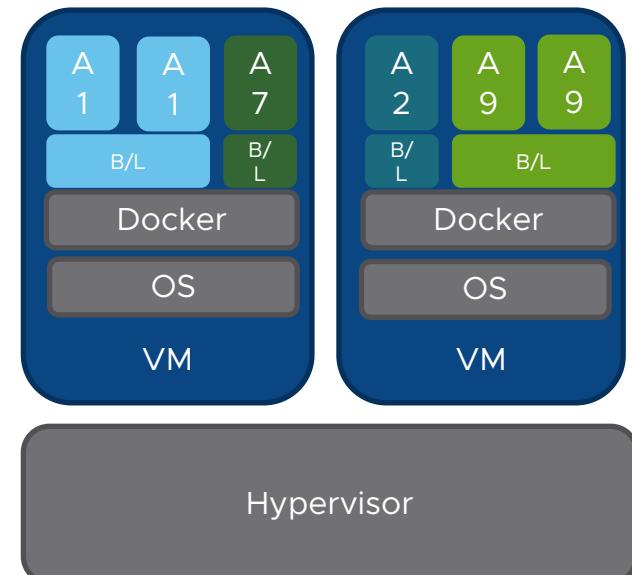


A container runs a process along with its dependencies

Share host OS kernel

Process isolation from other apps on same host

Lightweight – Just the app (NGINX, Redis, MySQL)

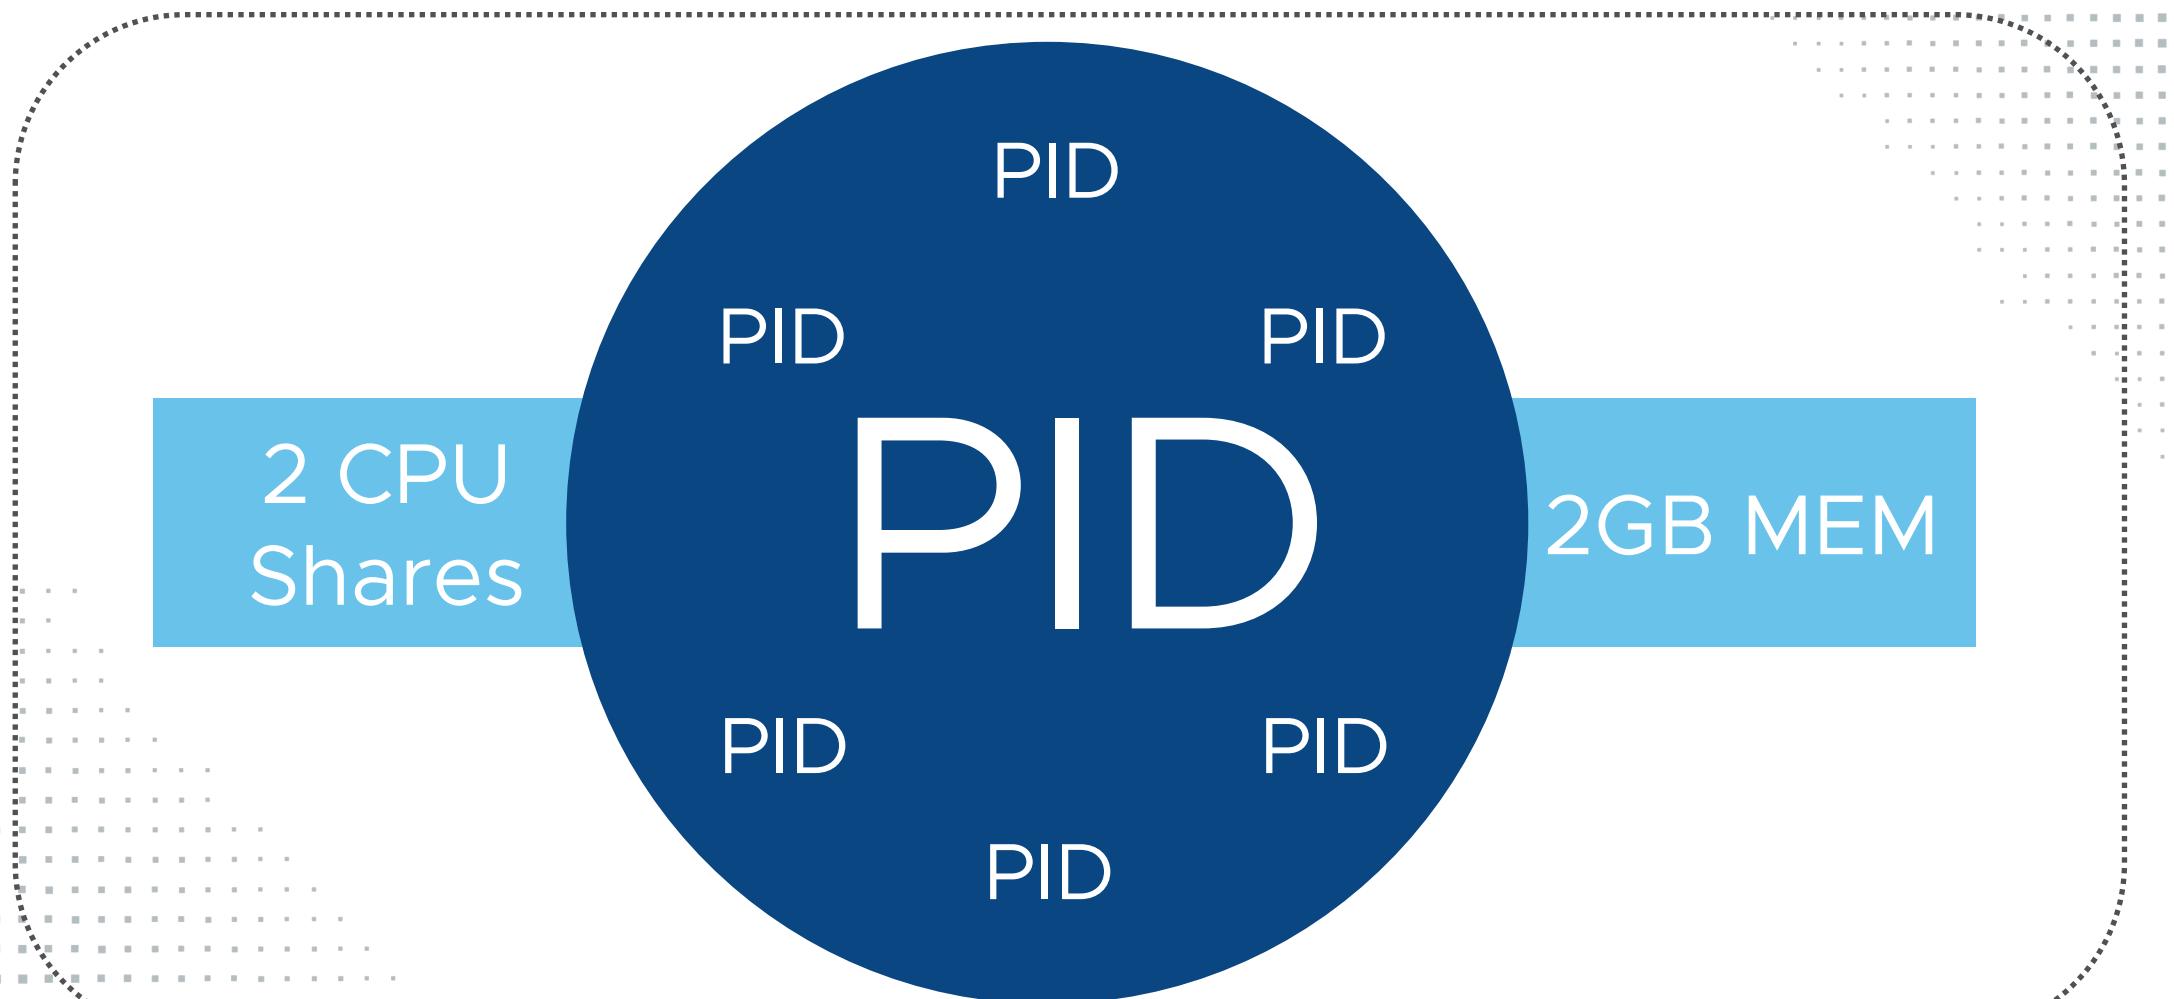


A scene from The Lord of the Rings: The Return of the King. Gandalf the White, with his long white beard and robes, sits in a large wooden chair, looking weary and sad. Saruman the White, with his pale face and dark robes, stands behind him, holding a staff. They are in a room with ornate wooden doors and a red tapestry in the background.

That

is a lie

How linux sees ‘containers’



Containers are a lie

What containers really are

```
vmware@k8s-node01a:~$ top
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
7714	root	20	0	604352	108412	37740	S	7.3	2.7	601:48.51	kubelet
1560	root	20	0	1188268	55248	8548	S	2.0	1.4	170:38.38	dockerd
24349	nobody	20	0	791068	285364	11412	S	1.7	7.1	101:51.55	prometheus
3951	root	20	0	45224	27868	17296	S	1.0	0.7	38:03.55	coredns
2690	root	20	0	43832	25944	16288	S	0.7	0.6	38:09.25	kube-proxy
24007	root	20	0	29132	7444	2892	S	5.9	0.2	9:40.25	alertmanager

```
vmware@k8s-node01a:~$ pstree -p 24349
prometheus(24349)--{prometheus}(24378)
| {prometheus}(24379)
| {prometheus}(24380)
| {prometheus}(24384)
| {prometheus}(24386)
| {prometheus}(24388)
| {prometheus}(24389)
```

```
vmware@k8s-node01a:~$ pstree -p 3951
coredns(3951)--{coredns}(3984)
| {coredns}(3985)
| {coredns}(3986)
| {coredns}(3987)
| {coredns}(3992)
| {coredns}(4009)
```

Control Groups

Allocating resources to a Namespace

```
vmware@k8s-node01a:~$ ls /sys/fs/cgroup/memory/kubepods/besteffort/pod0bbd0491-cbb9-11e8-8efc-005056a3e23d/  
18bcb52de19c8426afe024a5f364da352f65602a4d9ee5047c6a13ab59b1ae1c  memory.kmem.usage_in_bytes  
b203deb02da6774ba9a3f65a439a5b5135c1dc50436554350ce8608bb2276022  memory.limit_in_bytes  
cgroup.clone_children  
cgroup.event_control  
cgroup.procs  
memory.failcnt  
memory.force_empty  
memory.kmem.failcnt  
memory.kmem.limit_in_bytes  
memory.kmem.max_usage_in_bytes  
memory.kmem.slabinfo  
memory.kmem.tcp.failcnt  
memory.kmem.tcp.limit_in_bytes  
memory.kmem.tcp.max_usage_in_bytes  
memory.kmem.tcp.usage_in_bytes  
vmware@k8s-node01a:~$ cat /sys/fs/cgroup/memory/kubepods/besteffort/pod0bbd0491-cbb9-11e8-8efc-005056a3e23d/memory.kmem.limit_in_bytes  
640000000  
vmware@k8s-node01a:~$ ls /sys/fs/cgroup/cpuacct/kubepods/besteffort/pod0bbd0491-cbb9-11e8-8efc-005056a3e23d/  
18bcb52de19c8426afe024a5f364da352f65602a4d9ee5047c6a13ab59b1ae1c  cpuacct.stat      cpu.cfs_quota_us  tasks  
b203deb02da6774ba9a3f65a439a5b5135c1dc50436554350ce8608bb2276022  cpuacct.usage     cpu.shares  
cgroup.clone_children  
cgroup.procs  
cpu.cfs_period_us    notify_on_release  
vmware@k8s-node01a:~$ cat /sys/fs/cgroup/cpuacct/kubepods/besteffort/pod0bbd0491-cbb9-11e8-8efc-005056a3e23d/cpu.shares
```

Demo: Docker up

- 1. Pull basic hello world
 - 2. Validate image
 - 3. Run image
 - 4. Browse Website
 - 5. Attach to the namespace
-
- 1. Live demo on macOS
 - 2. Backup video

Kubernetes

Schedule at scale

Why is Kubernetes needed?

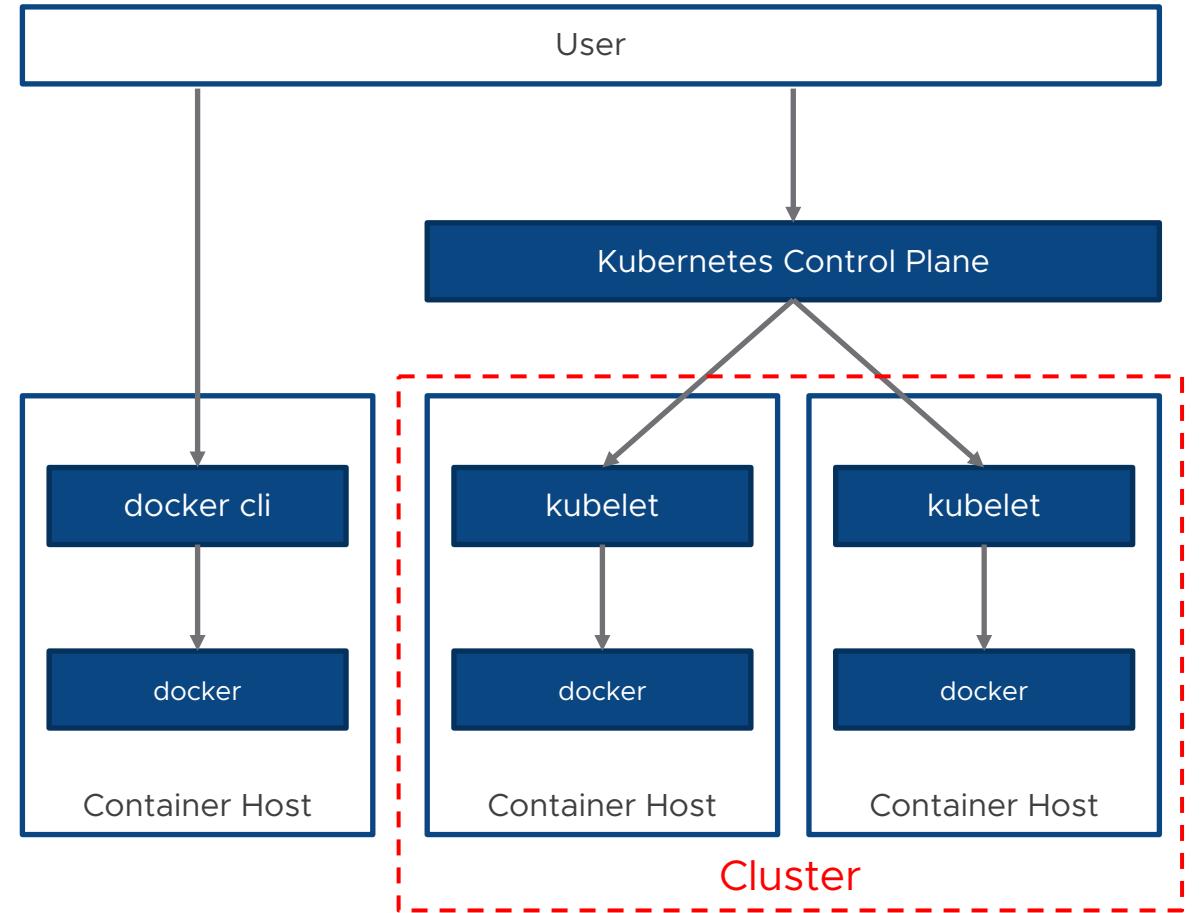


Orchestration across tens, hundreds, and thousands of containers

Developers don't care about storage/network/infrastructure nuances – they just want to deploy their application

Kubernetes centralizes orchestration providing

- Container Scheduling
 - AZs/Distribution/Behavior
- Container Management
 - Monitoring/Recovery/Updates/Logging
- Service Endpoints
 - Discovery/HA/Load Balancing/Auto Scale
- External Services

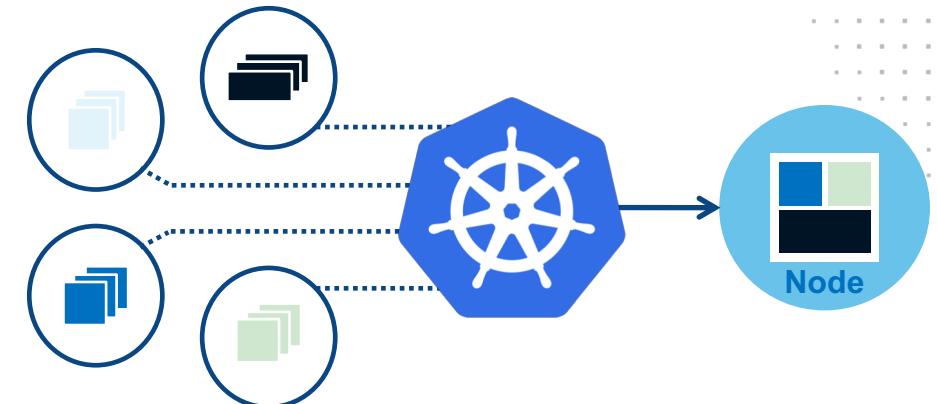


What is Kubernetes used for?

Kubernetes is an open-source system for automating **deployment**, **scaling** and **management** of containerized applications.

Major Features

- **Deploy** applications quickly and predictably
- **Scale** applications on the fly
- **Roll out** new features **seamlessly**
- **Optimize** use of hardware
- **Self-healing**(**Restart**, **replace** and **reschedule** containers)



Kubernetes

Basic Objects

Namespaces

ReplicaSets/ReplicationController

Pod

Service

Ingress

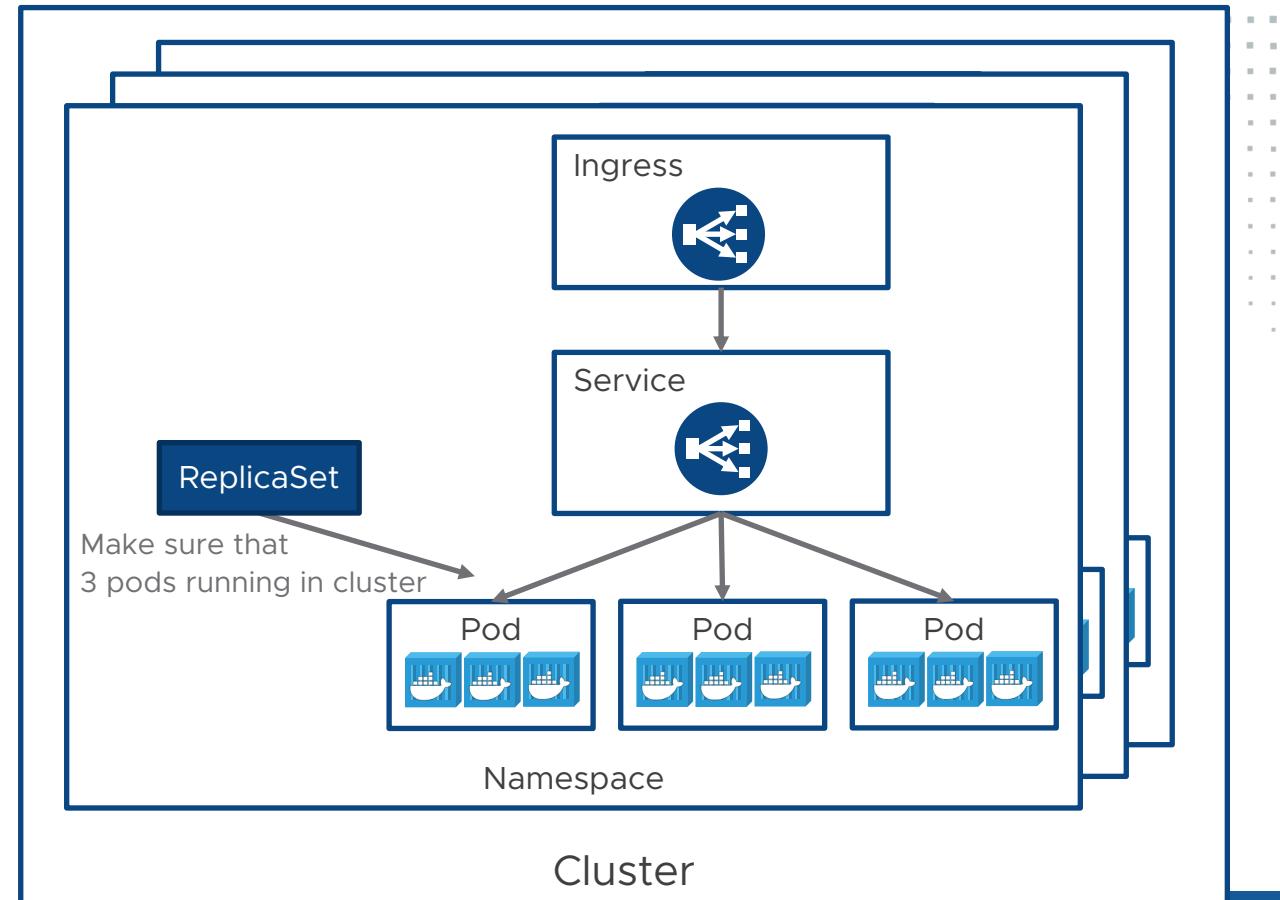
NetworkPolicy

Label

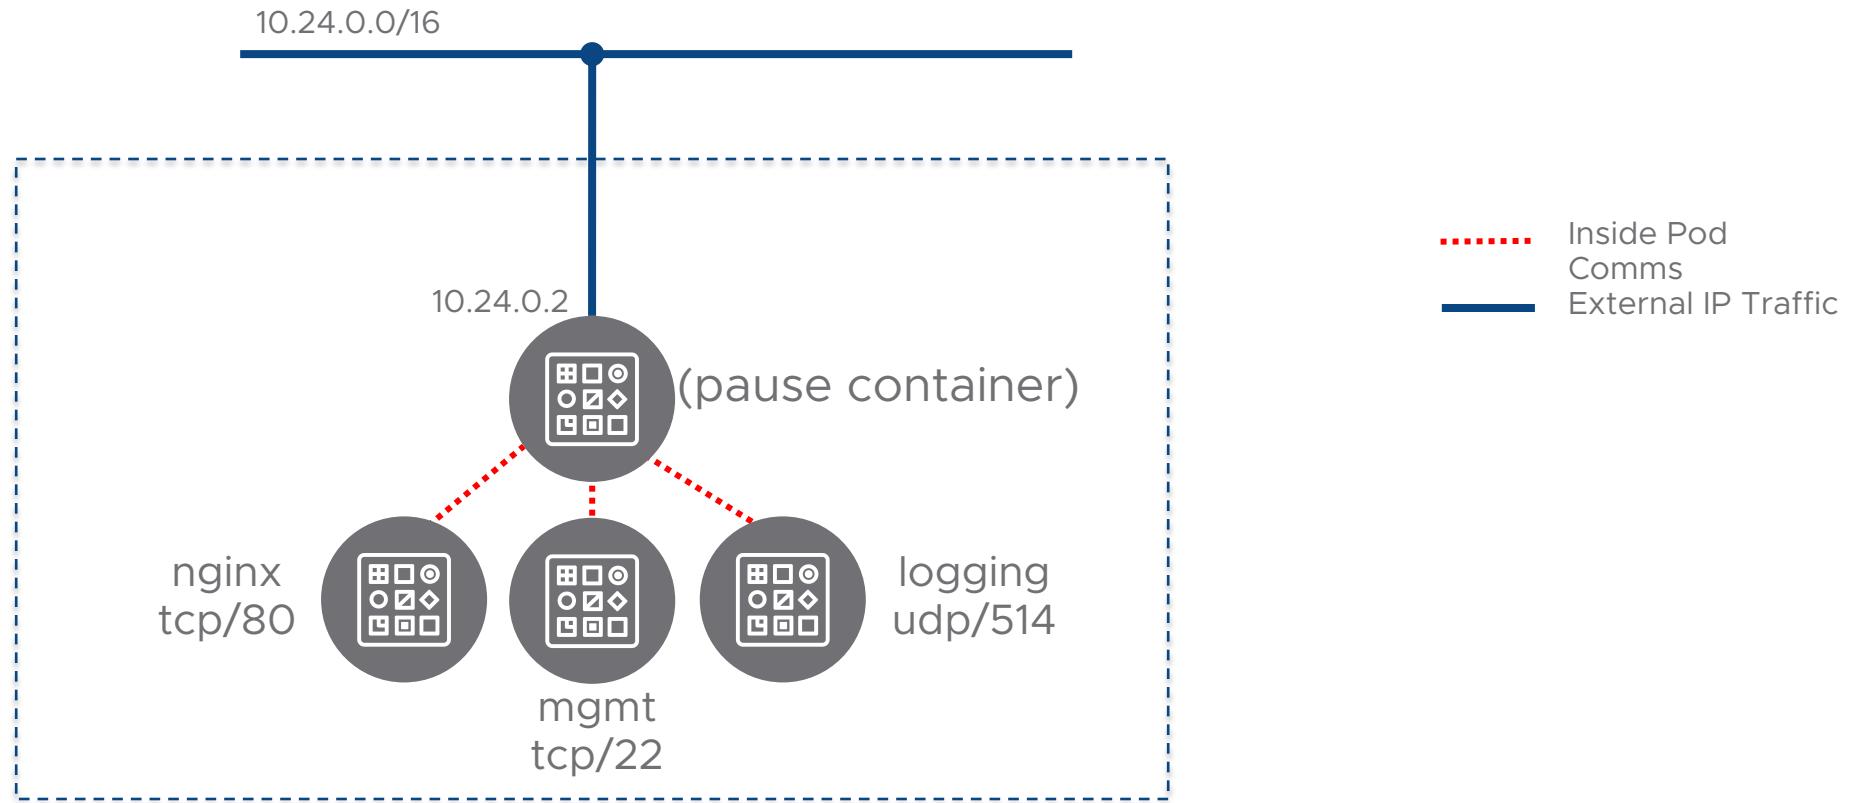
Kubernetes Cluster specific

Node
Cluster

vmware®



Kubernetes Pod



- A Pod is a group of one or more containers that shares an IP address and a data volume
- A function can be made up on many pods delivering a service such as nginx delivering a web function
- A namespace can have many pods providing different functions such as nginx, redis, mysql for apps.

Kubernetes Labels

Labels are used as metadata for a resource

Labels provide a way of matching for use in memberships, inclusion, exclusion, or identification

```
vmware@k8s-node01a:~$ cat example-svc.yaml
apiVersion: v1
kind: Service
metadata:
labels:
  app: nginx
  role: web
  name: nginx-demo-svc
spec:
  ports:
  - name: tcp
    port: 80
    protocol: TCP
    targetPort: 80
  selector:
    app: nginx
    role: web
  type: ClusterIP
```

In the case of multiple requirements, all must be satisfied so the comma separator acts as a logical AND (&&) operator.

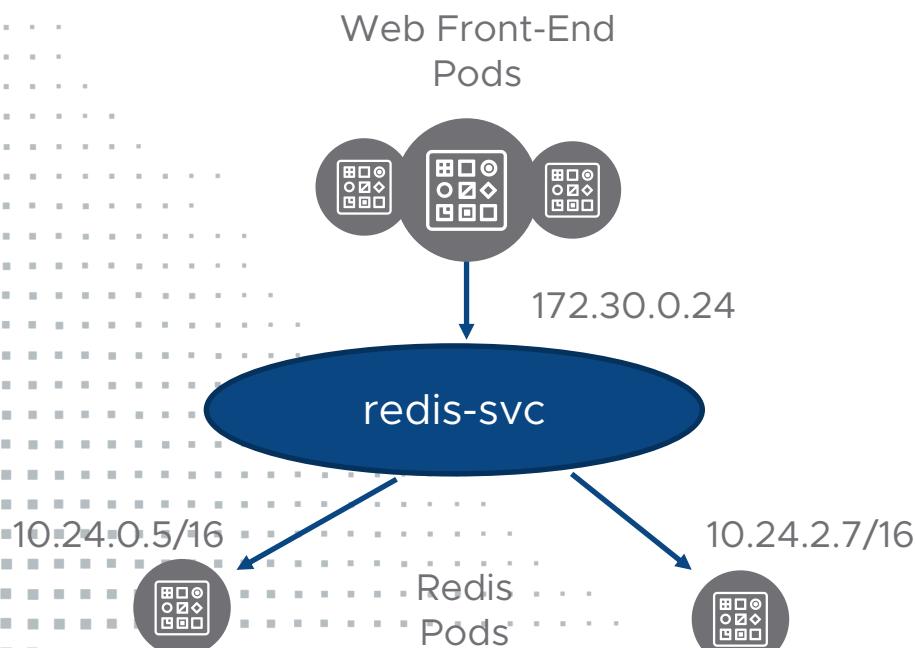
An empty label selector (that is, one with zero requirements) selects every object in the collection.

A null label selector (which is only possible for optional selector fields) selects no objects.

```
vmware@k8s-node01a:~$ cat example-networkpolicy.yaml
# Relatively Newer object supports matchLabels and matchExpressions
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: nginx-app-policy
spec:
  podSelector:
    matchLabels:
      app: nginx
      role: app
  policyTypes:
  - Ingress
  ingress:
  - from:
    - podSelector:
      matchLabels:
        app: nginx
      matchExpressions:
      - {key: role, operator: In, values: [web, db, client]}
```

Kubernetes Service

```
vmware@k8s-node01a:~$ kubectl describe svc redis-svc
Name:           redist-svc
Namespace:      default
Labels:         name=redis-svc
Selector:       name=redis-svc
Type:          ClusterIP
IP:            172.30.0.24
Port:          <unnamed> 6379/TCP
Endpoints:    10.24.0.5:6379, 10.24.2.7:6379
```



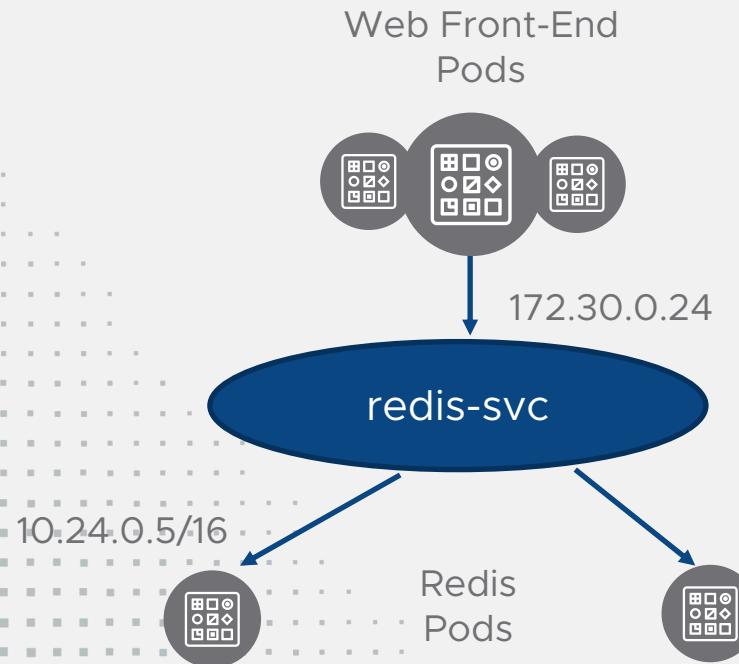
vmware®

Kubernetes Service

- **Gist:** A Kubernetes Service is an abstraction which defines a logical set of Pods
- **East/West Load-Balancing:** In terms of networking a service usually contains a cluster IP, which is used as a Virtual IP reachable internally on all Nodes
- **IPTables:** In the [default](#) upstream implementation IPTables is used to implement distributed east/west load-balancing
- **DNS:** A service is also represented with a DNS names, e.g. '*redis-svc.cluster.local*' in the Kubernetes dynamic DNS service (CoreDNS) or through environment variable injection
- **External Access:** A K8s Service can also be made externally reachable through [all](#) Nodes IP interface using 'NodePort' exposing the Service through a specific UDP/TCP Port
- **Type:** In addition to ClusterIP and NodePort, some cloud providers like GCE & OpenStack support using the type 'LoadBalancer' to configure an external LoadBalancer to point to the Endpoints (Pods)

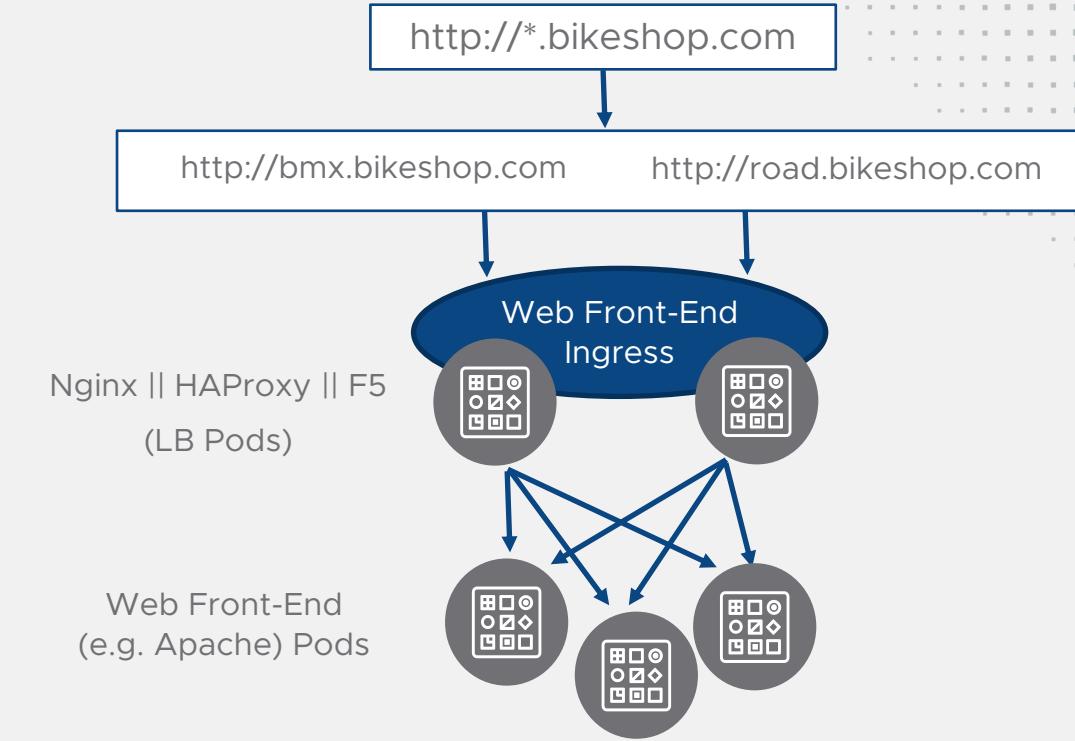
Kubernetes Load Balancing

East-West Load Balancing



East-West Load Balancing is provided through Kubernetes Service using ClusterIP & iptables

North-South Load Balancing



Can be achieved through Kubernetes Ingress or External third Party Load Balancer using NodePort

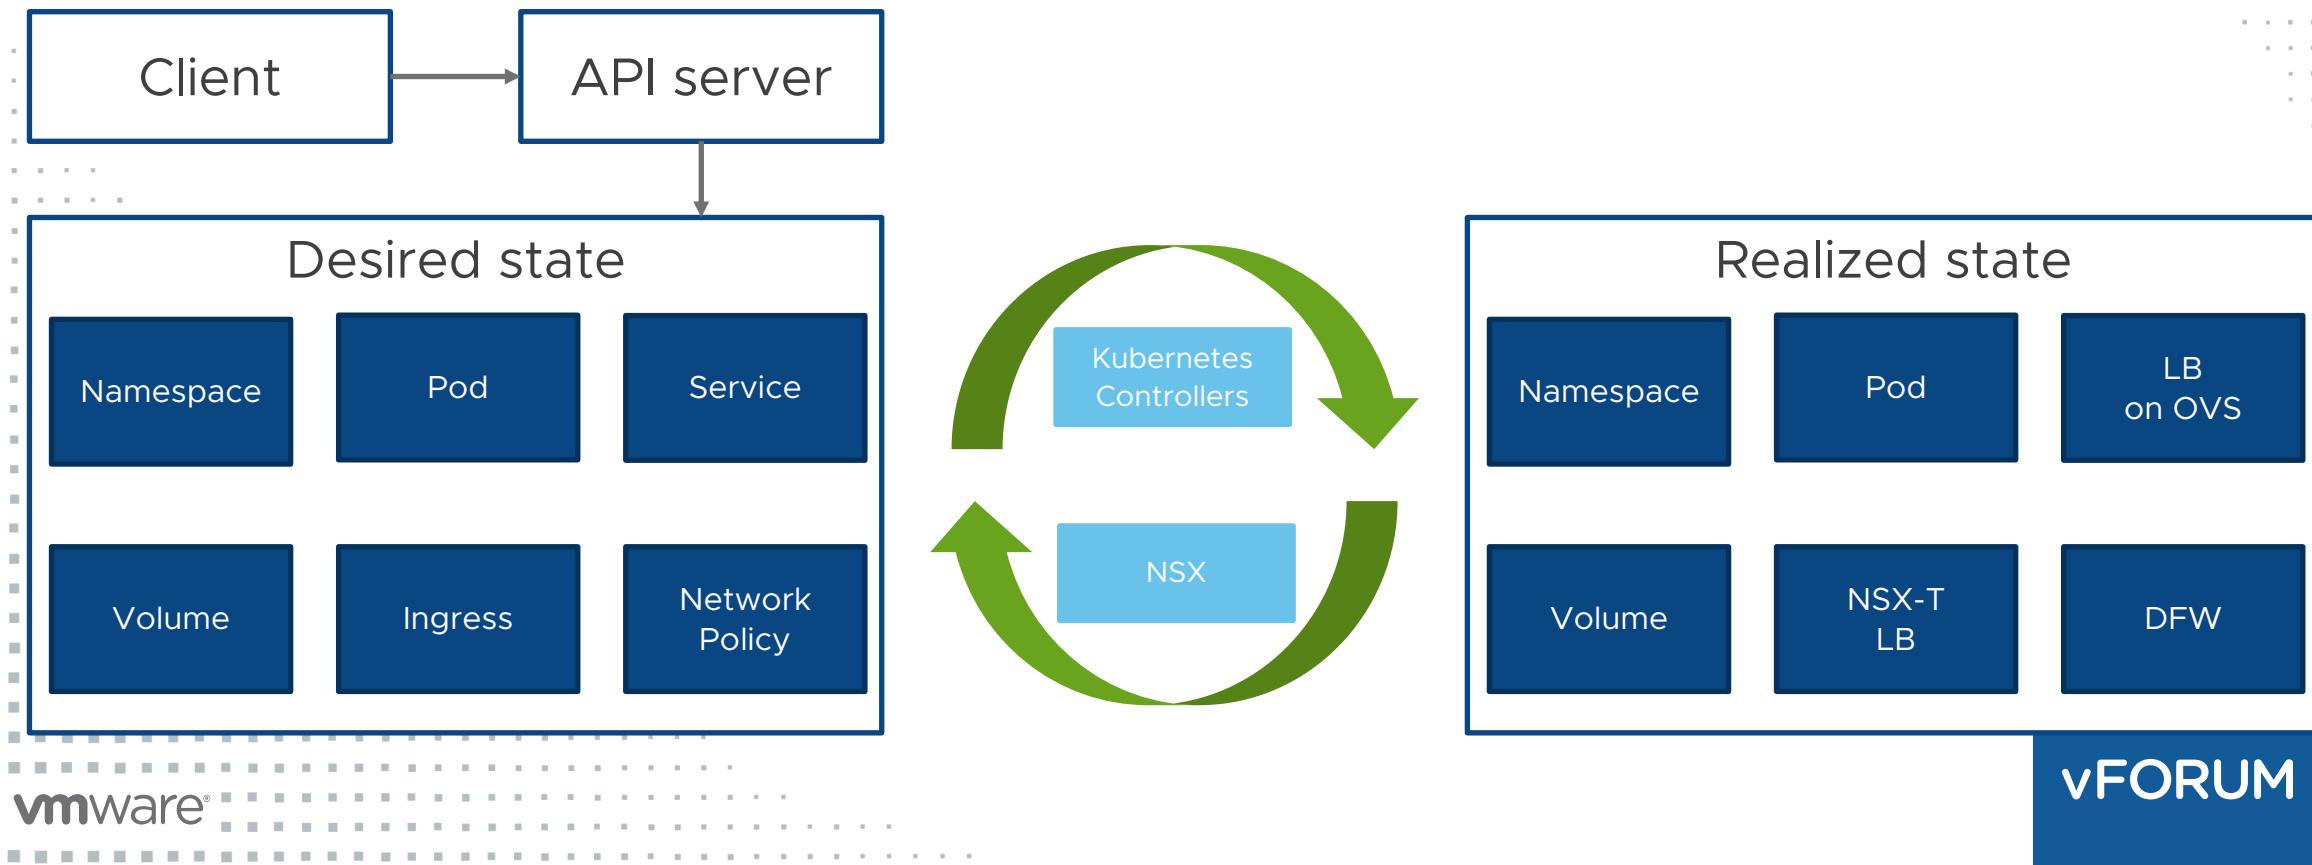
Kubernetes Objects (Resource Types)

Component	Description
Pods	A grouping of one or more containers as an atomic unit
Namespaces	A way to organize items in a cluster (+ RBAC)
Labels, Annotations & Selectors	Tags for component grouping and methods to access them
Service Discovery	An object associated to a label selector to provide a LB and Service DNS
ReplicaSets	A cluster wide Pod manager providing Pod scaling
DaemonSets	A Pod manager to ensure a Pod is scheduled across a Cluster Node set
StatefulSets	Replicated Pods where each Pod gets an indexed hostname
Jobs	A Pod which runs until the process returns a successful termination
Deployments	Manages the release rollout of new versions of Pods

How does Kubernetes do its thing?

Concept

Kubernetes control plane work to make the actual state based of desired state which is made by user/client. There are some running controllers which has loop which gets rid of loop between desired state and realized state.



Demo: Kubernetes

- 1. Schedule a single nsx-demo RC
- 2. Scale pods
- 1. Live demo via Singapore
- 2. Backup video

Kubernetes and NSX-T

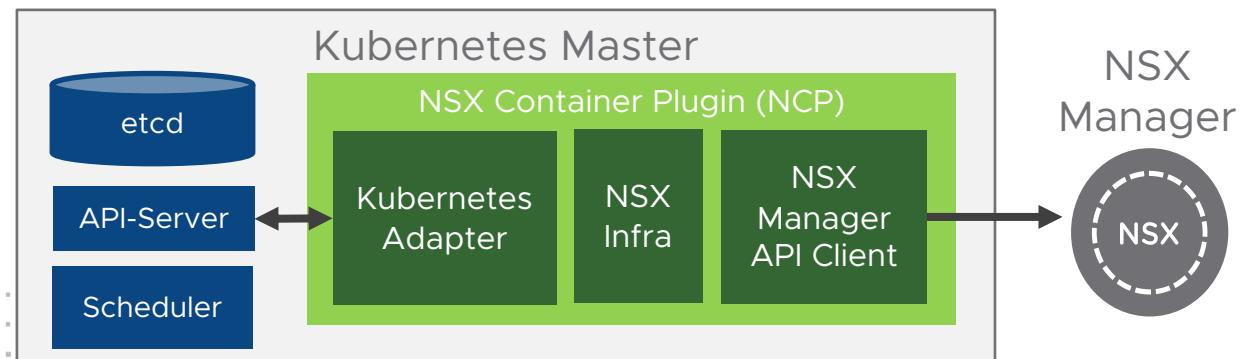
NSX-T as a Network Provider

NSX-T Container Plugin

Facilitating the NSX-T - K8s Integration

NSX Container Plugin (NCP) is the core component that provides integration between K8s and the NSX Manager.

- Runs as a [container](#) inside a K8s pod
- [Monitors for changes](#) of relevant objects (namespaces, pods) on K8s API server
- [Creates the relevant NSX Objects](#) (logical switches, routers) on creation of K8s constructs
- NCP is built in a [modular way](#), so that individual adapters can be added for different CaaS and PaaS systems

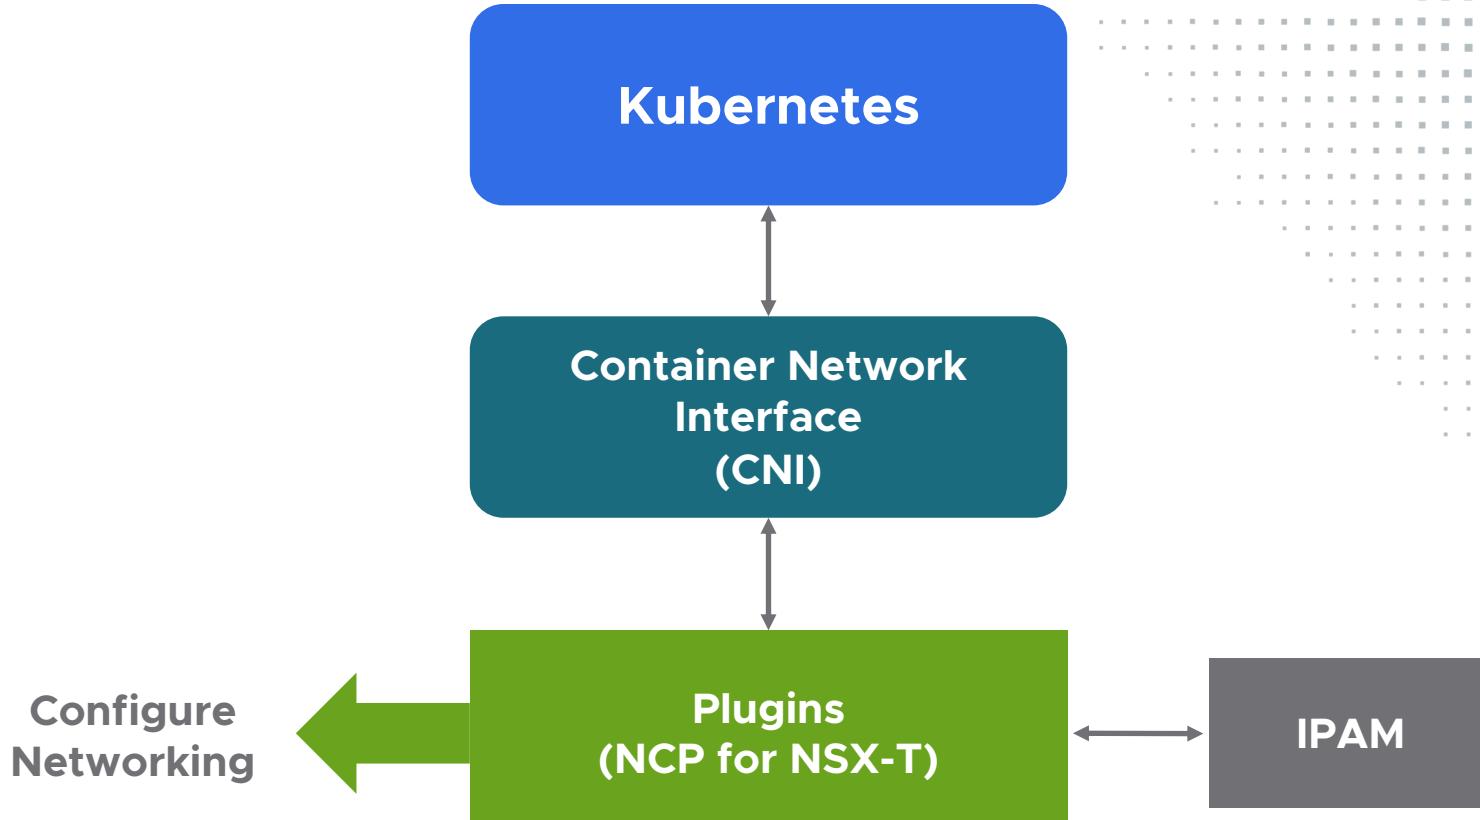


Container Networking Interface

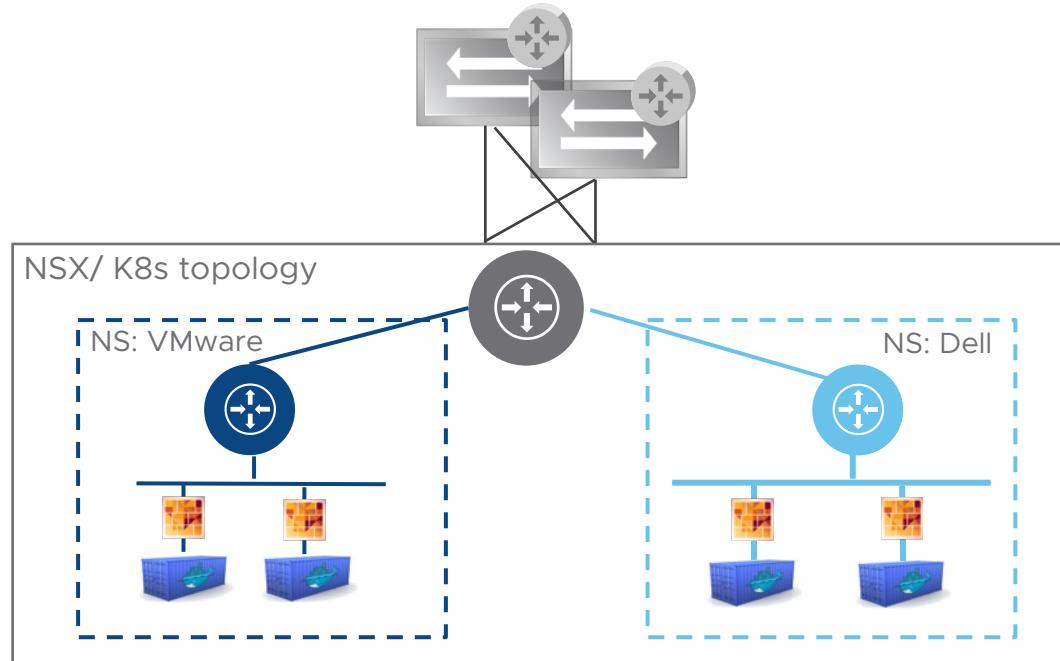
Connecting the Dots ...

Network Container Plugin (NCP)

- Common interface between container runtime and implementation
- Plugins responsible for the network plumbing – IPAM/switching/firewall/load balancing – for container workloads
- Kubernetes commands stay the same; NSX does networking functions
- NSX Container Plugin (NCP) results in the NSX-T to Kubernetes integration



Kubernetes NSX Topology

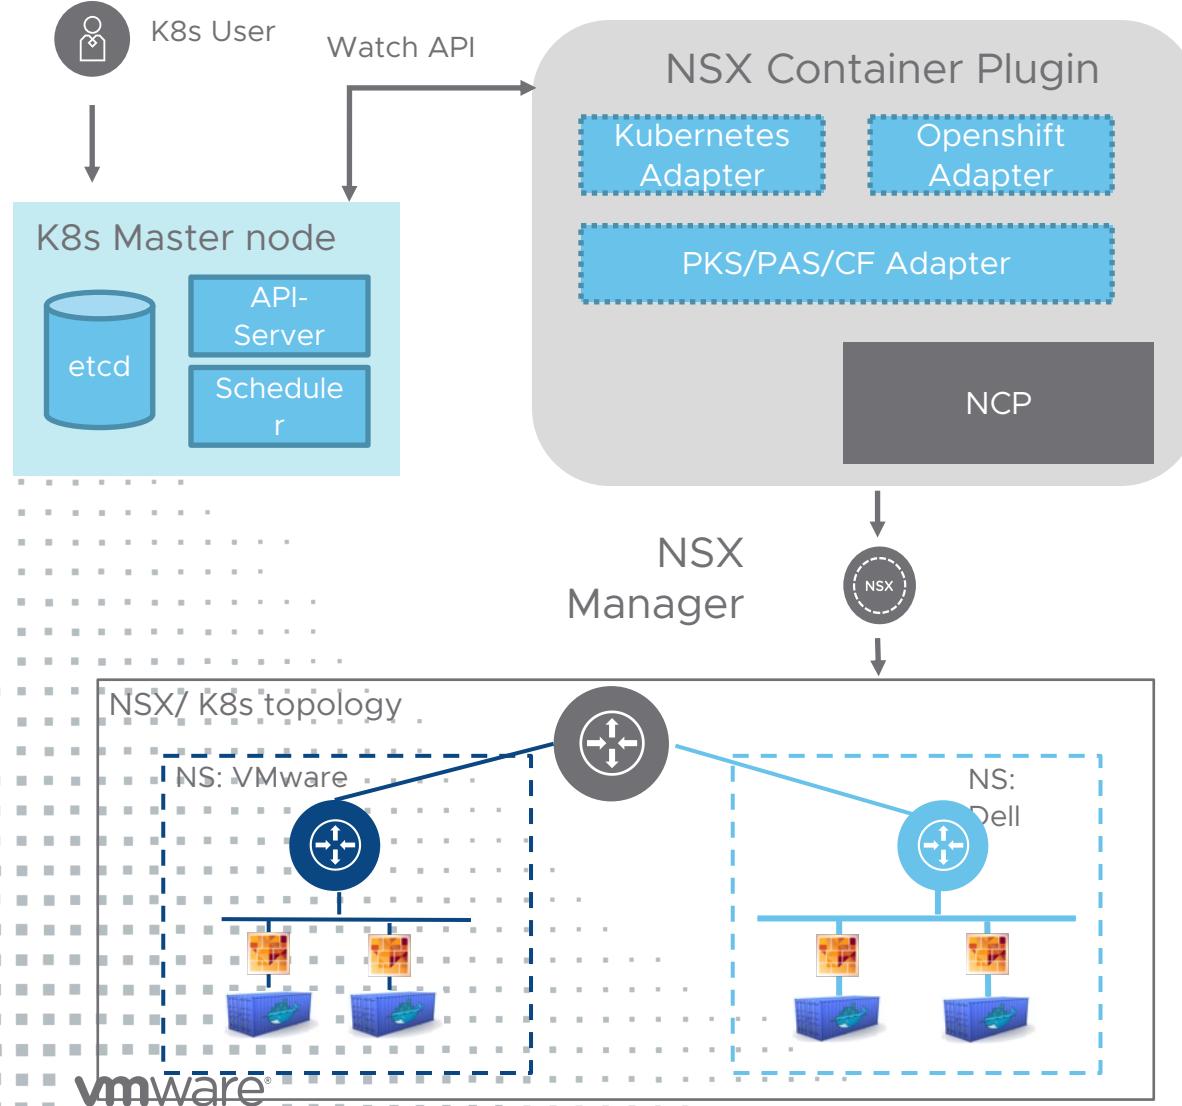


Kubernetes NSX Topology

- **Namespaces:** Dynamically builds a separate network topology per K8s namespace, every K8s namespace gets one or more logical switches and one Tier-1 router
- **Nodes:** Are not doing IP routing, every Pod has its own logical port on a NSX logical switch. Every Node can have Pods from different Namespaces and with this from different IP Subnets / Topologies
- **Firewall:** Every Pod has DFW rules applied on its Interface
- **Routing:** High performant East/West and North/South routing using NSX's routing infrastructure, including dynamic routing to physical network
- **Visibility and troubleshooting:** Every Pod has a logical port on the logical switch with:
 - Counters
 - SPAN / Remote mirroring
 - IPFIX export
 - Traceflow / Port-Connection tool
 - Spoofguard
- **IPAM:** NSX is used to provide IP Address Management by supplying Subnets from IP Block to Namespaces, and Individual IPs and MAC to Pods

Kubernetes and NSX components

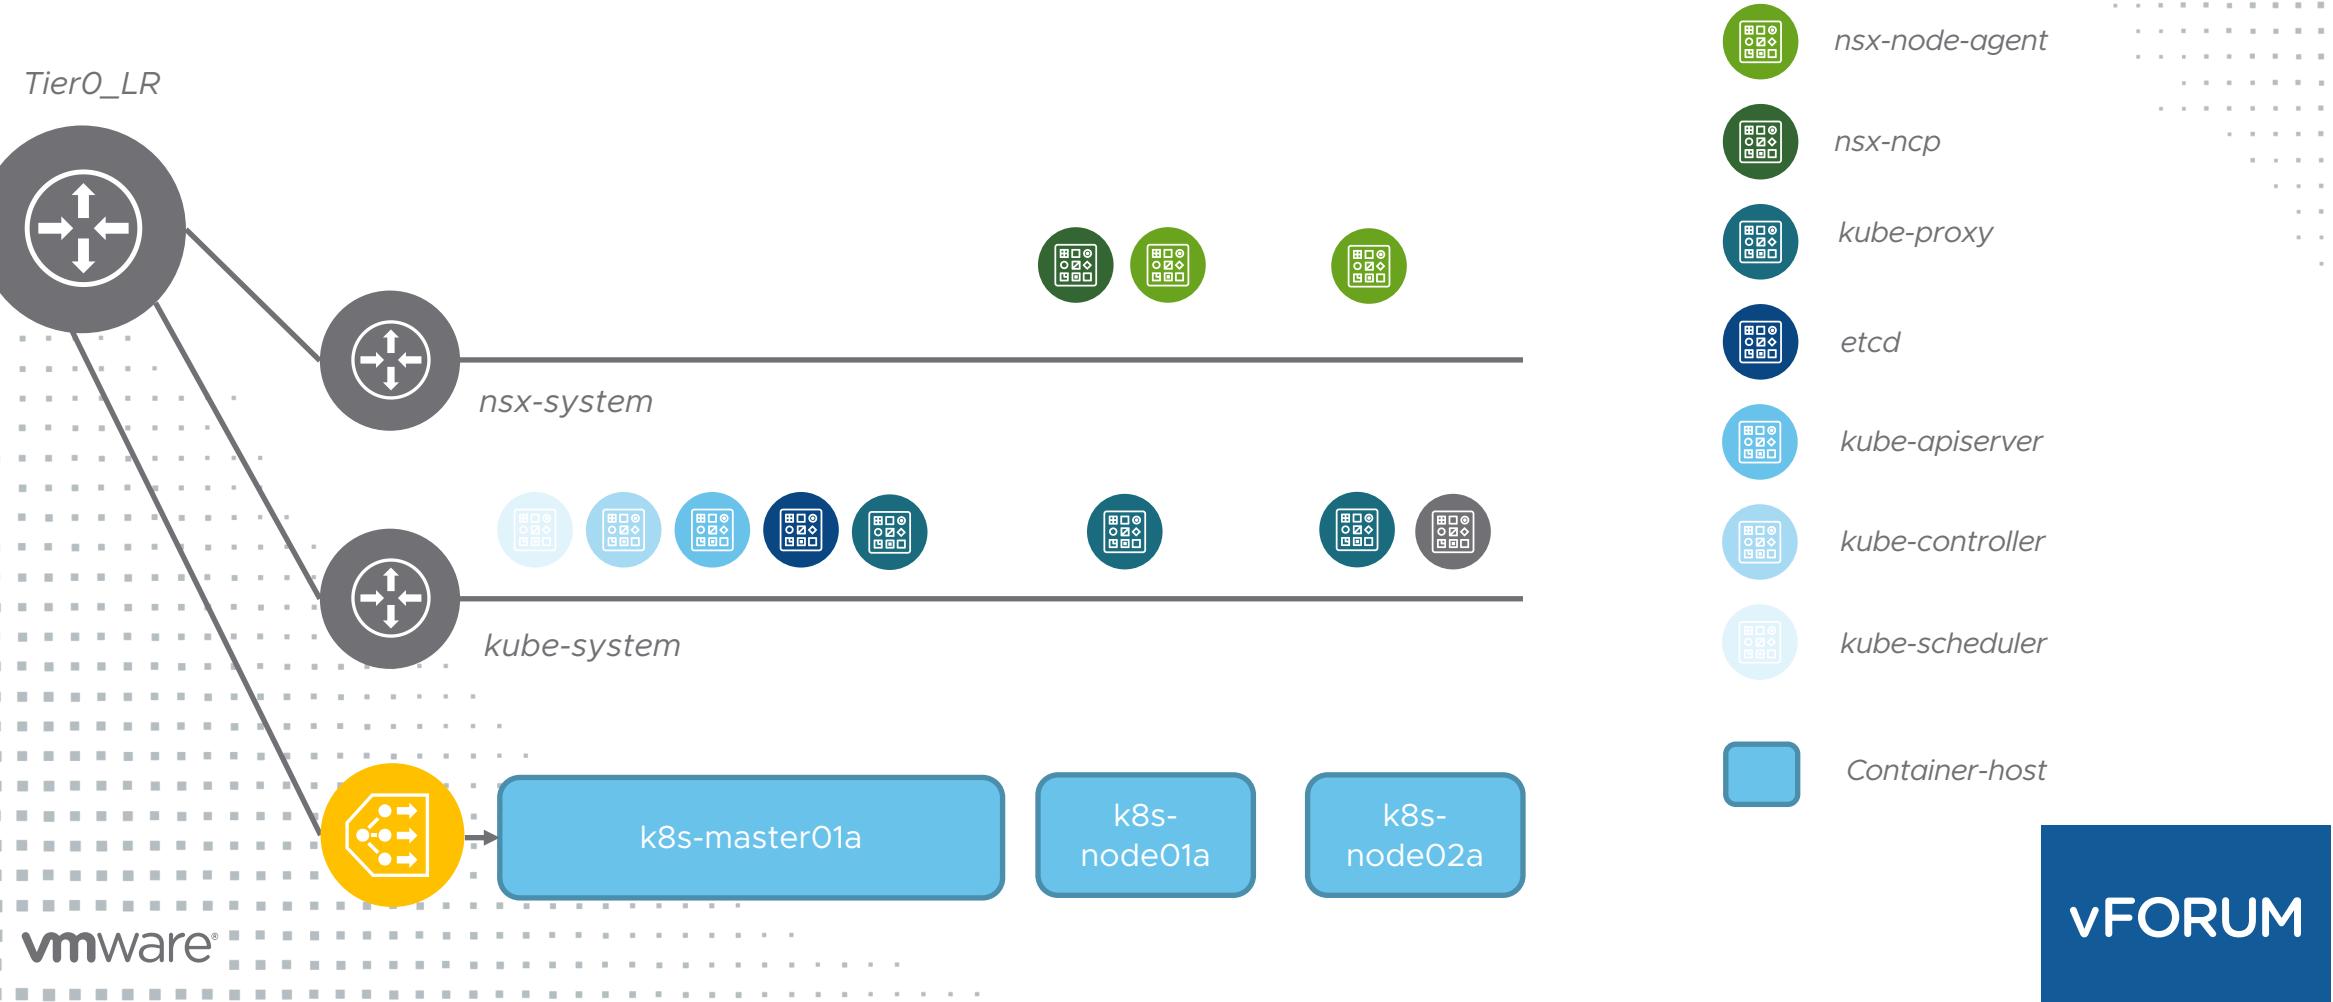
Network Container Plugin (NCP)



Network Container Plugin (NCP)

- **NSX Container Plugin**: NCP is a software component provided by VMware in form of a container image, e.g. to be run as a K8s Pod
- **Adapter layer**: NCP is build in a modular way, so that individual adapters can be added for different CaaS and PaaS systems
- **NSX Infra layer**: Implements the logic that creates topologies, attaches logical ports, etc. based on triggers from the Adapter layer
- **NSX API Client**: Implements a standardized interface to the NSX API

Ground floor topology for NSX-T and Kubernetes integration



Planespotter

Putting it all together



What is Planespotters?

Planespotter is a free, open-source application,
built with a “Cloud-native” mindset

- Built for demonstration of key Kubernetes and NSX resources
- Pods, Replication controllers, StatefulSets, NetworkPolicy, Ingress, Services

Free!

- You can get planespotter
- You can use planespotter
- You can use planespotter on Minikube, GKE, VKE, AKS, Vanilla K8s.

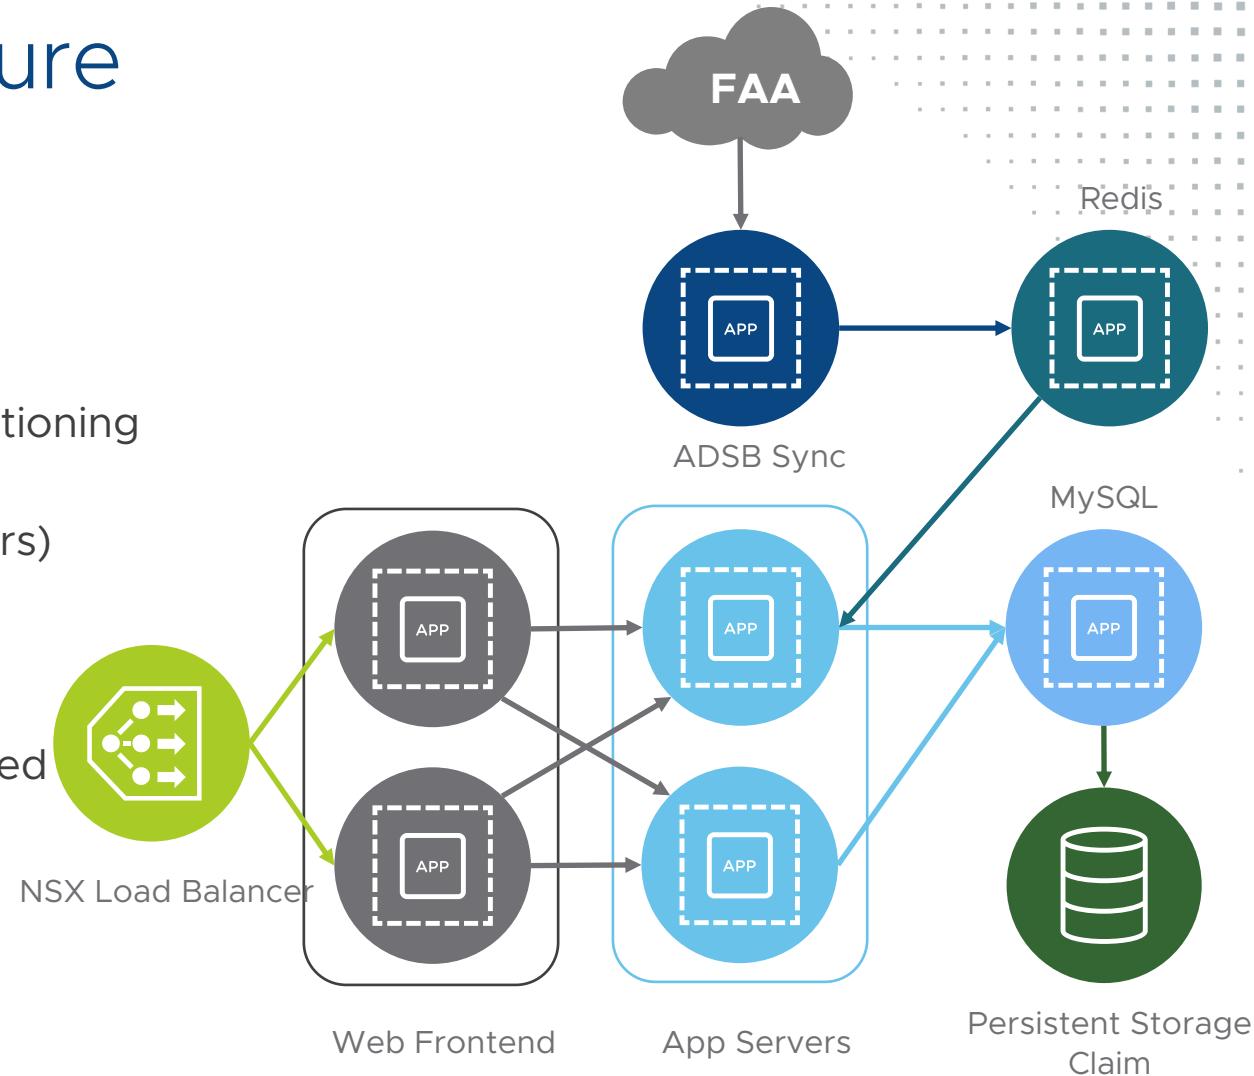
Planespotter architecture

Defined by 5 YAML files

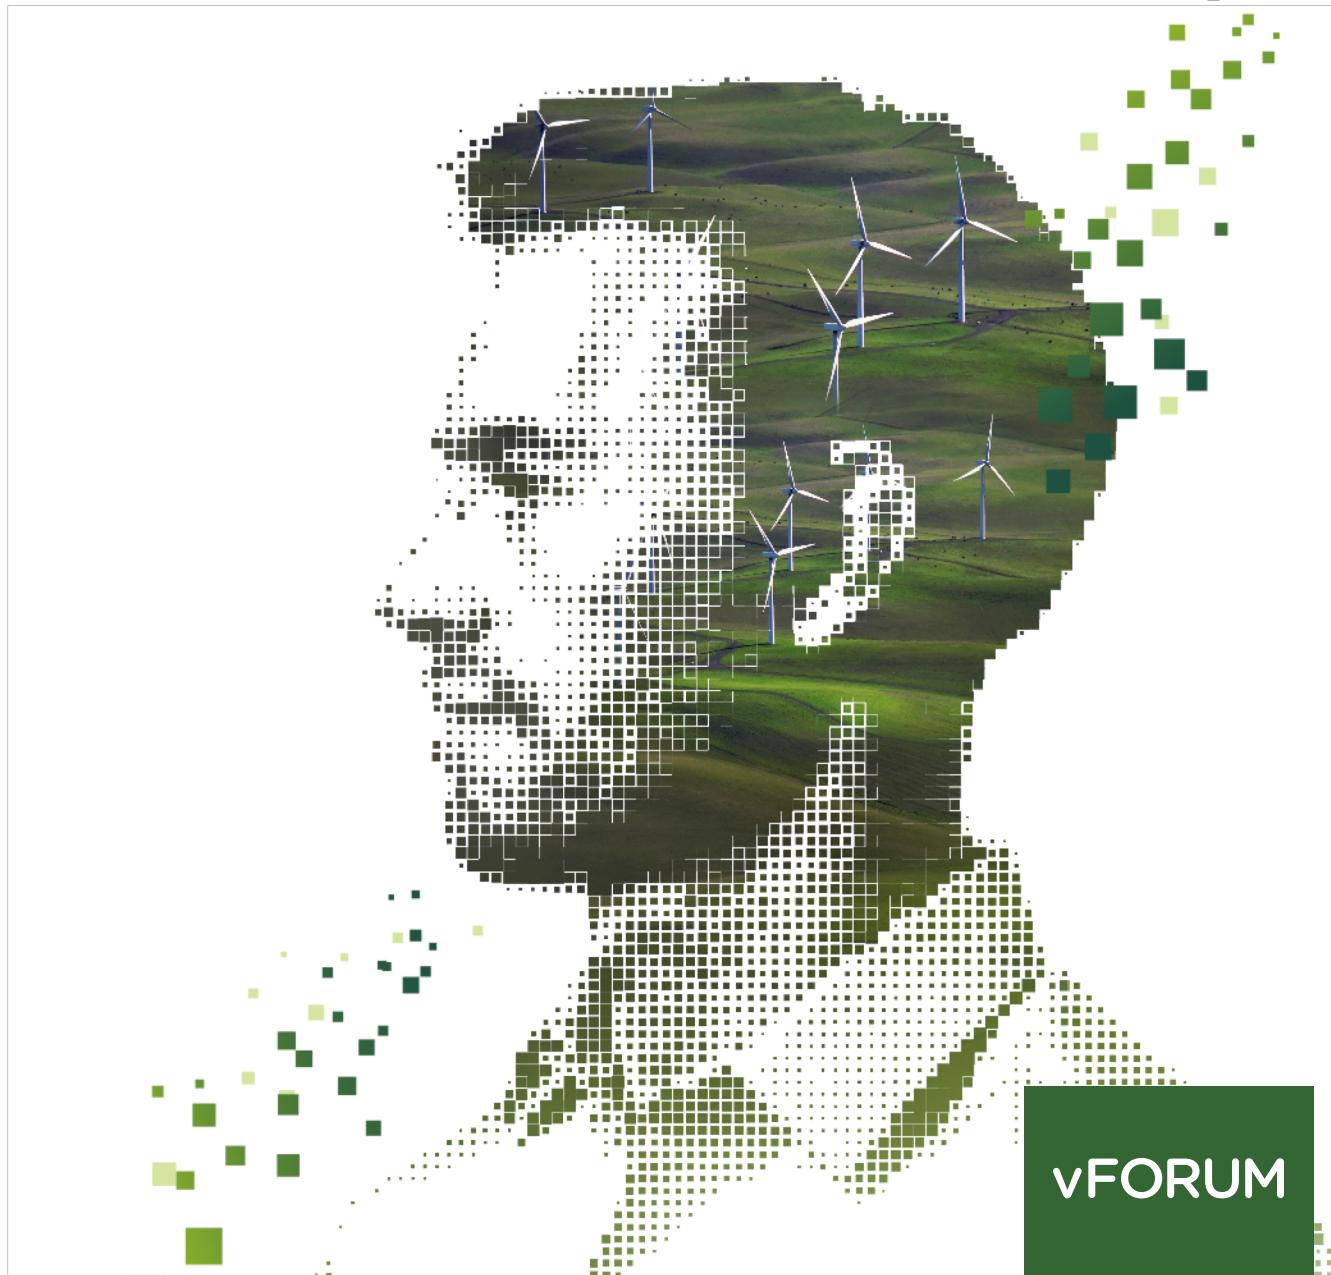
- Mysql – storing static plane information
- Redis - in-memory DB for real time positioning
- Live Position ingestion
- API and Correlation servers (App Servers)
- Web Frontend

Pre-reqs

- Kubernetes/NSX-T integration configured
- Storage class defined or a default set
- Any ops tools configured



DEMO TIME



Putting Planespottter together

Demo time

The order of operations to create the planespotter app is as follows:

- Create Namespace
- Create a Storage Class
- Create Persistent Volume Claim
- Create a Stateful set
- Create an application deployment
- Create a frontend deployment
- Create a redist deployment
- Apply network policies

Use Kubernetes to validate
Use NSX-T ops tools to validate

Coming up next

Migration, Connectivity and Security for the Hybrid cloud with NSX Hybrid Connect

1:30PM – 2:30PM

Thank you

**POSSIBLE BEGINS
WITH YOU**



vFORUM